

Q2: Bads and blames	1
Intro	2
Test cases	2
Data hoarder	3
Argument tweeter	3
Internet isolation	3
People Ratings Agencies	4
Harms of informational privacy violation	4
Methodology	4
The problem with starting from privacy	5
Focus on harmful violations of privacy	5
The concept of harm	6
Harm principle	6
Definition of harm	7
Setback of interest	7
Interests	7
Well-being	8
Setbacks	9
Death	9
Other problems	10
Wrongfulness	10
Odd cases	11
Harmless rights violations	11
Hurting without harming	12
Informational harms -- setbacks	13
Exploiting weaknesses	13
Stalking and worse	14
Threats to autonomy	16
Relationships	17
Democracy	17
The panoptic sort	??

Identity theft.....	??
Informational harms -- rights violations.....	17
What kind of rights?.....	18
Non privacy rights.....	18
Right to protected sphere.....	18
Accounts of right to protected sphere.....	19
Negative Liberty account.....	19
Positive liberty.....	20
Negative liberty.....	21
Right to protected sphere.....	21
Objections.....	22
Republican account.....	22
Right to protected sphere.....	23
Objections.....	23
Relational account.....	24
Right to protected sphere.....	24
Objections.....	25
Can machines violate privacy?.....	25
No. Understanding matters.....	25
Application of 3 approaches.....	26

## 1 Intro

We've discussed when an ethical company may profit from the use of personal data.

Let's now turn to when things go wrong and take up a different question:

(Q2) How should we assign blame when people are harmed by algorithmic uses of personal data?

This discussion will require us to first get the hang of how people could be harmed in this way and then discuss how our ideas of blame and responsibility may need revision in light of new technology.

## 2 Test cases

As we work through discussions of harm and blame, it will help to keep in mind some made up cases in which things go wrong with personal data.

Your mileage may vary. Some of these may seem fairly innocuous. Some may seem creepy but not necessarily harmful. Others may seem straight-up despicable. Take a minute as you go through them to try to put your finger on what is wrong about the ones that seem wrong. More importantly, pay attention to tweaks to the details that flip your reaction from ‘meh’ to ‘oh hell no’. Those are exactly the things we’re going to try to uncover in understanding how uses of personal data can harm.

## **2.1 Data hoarder**

*Data hoarder:* Red is a data hoarder. She has an intense and obsessive desire to possess data. However, she doesn’t actually care about the contents of the data. She just likes to know that she possesses it. Suppose —contrary to the real world — that she keeps her data in a perfectly secure database which only she can access. She doesn’t look in the database; she just cares about having data, not doing anything with it. Indeed, we know that she’ll never use it in any way which affects people’s choices.

## **2.2 Argument tweeter**

*Argument Tweeter:* A new app which runs on the various ‘smart’ devices in your home —your Alexa, iPhone, et cetera — detects any time there is an argument in its vicinity. The recording is immediately tweeted out.

### Variations

There’s nothing identifiable about the clip

The twitter account is unpopular, no one ever hears the clip

Instead of being tweeted, it is stored in a database

Sex rather than arguments

Instead of tweeting them out it stores the audio files in a database [non-identifiably tagged; identifiably tagged]

It creates a hash (irreversibly encryption) of the audio files and blinks out the resulting hashes in morse code on an LED in the middle of the desert.

### Note

This raises issues of when the actual violation of privacy occurs. Is it when the clips are accessed? Or when they are collected.

What if it was set up in a way that you know your argumentative style will never be picked up

## **2.3 Internet isolation**

*Internet Isolation:* Suppose you, like many people, connect with your friends and family via the social media platform SnapFace. Due to an unforeseen confluence of parameters, the site's algorithm stops showing your posts to others. You don't know this. But you do notice a slide in likes / comments down to zero. At first it doesn't bother you. But, as time goes on and no one responds to even your best material —your most clever observations, your cutest outfits, the most adorable kitten/puppy/baby pictures— the perceived isolation and shunning start eating at you. You become socially anxious and depressed ..

## **2.4 People Ratings Agencies**

*People Ratings:* Panopticon LLC creates a web of partnerships and data-sharing agreements with every major data-collection company. Through these connections, Panopticon LLC has access to nearly everything each American has ever said or done online, as well as in the vicinity of a camera, and every purchase they've made.

Here are just a few of the revolutionary services they provide:

*Predictive Policing:* Panopticon offers a service to law enforcement which allows seamless access to the identities of every person within a few blocks of the scene of a crime. This list of people is served up along with a proprietary rating of the likelihood that each person would create a crime.

*Employee Screening:* Panopticon offers a service to employers with custom ratings of prospective employees in terms of who is likely to be insubordinate, who will not complain about menial work, and other work-related dispositions.

*Date Screening:* Through a new partnership with Tinder, Grinder, Bumble, and all the other major online dating apps, Panopticon offers ratings of prospective partners on everything from conscientiousness, cleanliness, likelihood of cheating, expected number of dates-before-sex, financial status, and other metrics. By offering users the ability to provide post-date/post-breakup ratings, it also provides ratings on kinkiness and sexual performance. [I'll stop here; this is creepy enough]

If these services do not seem problematic to you, feel free to add new details until they strike you as worrisome. All set? Good. Keep that in mind in what follows.

## **3 Harms of informational privacy violation**

### **3.1 Methodology**

Thus far I've tried to avoid discussing things in terms of privacy. But thinking through our cases, that's no longer going to be avoidable. When we ask what's going wrong in these cases, the irresistible answer is that people's privacy is being violated.

Hopefully at the point in class you're expecting us to dive into getting clear about what privacy is before applying the concept. If so, yay! This is usually a great strategy. Too many confusions and pointless disagreements arise when people think they're arguing about the same thing but are actually talking about different things.

### **3.1.1 The problem with starting from privacy**

But I'm not sure that starting by clarifying the concept of privacy is the best approach here. That's in part because the issues we're dealing with are likely to some extent genuinely new. Thus existing accounts of privacy may or may not be up to the job of dealing with them.<sup>1</sup>

More importantly, our task here is a bit different from what we've been doing elsewhere. Previously, we dealt with cases where we weren't really sure what's right and wrong. Think back to Holley. We spent a lot of time trying to come up with some principle that could help us determine when it's okay for a salesperson not to divulge information. A lot of the cases we discussed left a lot of room for reasonable disagreement about what's right and wrong.

The way I've constructed our task here is different. We're starting with cases which nearly everyone will agree involve unacceptable activities. If the cases I've come up with don't elicit those intuitions for you, I'd bet that we could tweak them until you'd agree that they involve something wrong.

If we're talking about cases which all involve something wrong, and bring in a notion of privacy which can't capture the wrongness, that's a strike against the account of privacy. It doesn't show us anything about how to understand what's wrong in all the cases. Put a different way, no account of privacy is likely to change our minds about the wrongness of these cases. That's why I want to avoid talking about the concept of privacy as much as possible — it's not likely to help us in this task.

---

1. Obviously, to really establish that we'd need to work through all of the existing accounts and carefully consider their strengths and weaknesses for dealing with our concerns. That's a Ph.D. dissertation; we're not going to do that here.

### **3.1.2 Focus on harmful violations of privacy**

So, what's the alternative approach? Well, notice that sometimes we might not care very much when privacy is violated. Think of someone accidentally walking in on you in a store changing room; not great, but usually not a huge deal. But we do care a lot when people suffer harms from their privacy is violated.

Thus the alternative is to start by thinking about how to explain how privacy violations can be harms. Approaching in this way, will also give us some resources to draw on. In particular, we can frame the discussion in terms of a well-worked out concept of harm .

If this approach still seems suspect, here's another analogy. Legal scholars sometimes discuss what are called *harmless trespasses*. If you own some land and I take a shortcut across it, leaving no trace, I have still violated your legal right to control who enters your property. But I think most people would agree that prosecuting someone for a harmless trespass should normally be at the very bottom of a prosecutors' priorities. We think the prosecutor should focus first on trespasses that involve harms (e.g., where I chop down your trees and knock down your fence dragging them away).

## **3.2 The concept of harm**

We are going to start with the pretty clear intuition that people are harmed in our cases. This allows us to get some leverage on our problems since harm is a much better worked out concept, at least in some areas of philosophy.

I'm thus going to draw on an influential account of harm from philosophy of law that comes from work by Judy Thompson and Joel Feinberg. Like every view it has problems, some of which I'll mention as we go, but we need not focus on those here.

A bit of warning in advance, I'm going to go into a lot of detail about the concept of harm in what follows. We will need those details when we turn to applying it concept. But don't worry if you don't follow all of them on the first read of this. They are more here for reference so that you can refer back to them when you need them.

### **3.2.1 Harm principle**

To begin, I'm going to stipulate that, on our use, harms and acts which cause harm are always grounds for public action or complaint. If you've ever heard of the 'harm principle' which comes up in discussions of criminal law, that's what's in the background here. The harm principle says something like:

(HP) An activity may be criminalized only if it causes harm to others

For those of you who have heard of the harm principle, that's probably because of the fact that this is a necessary condition. If we accept HP, it follows that actions which do not harm other people cannot be criminalized. Thus, at least in spirit, it comes up often in discussions of drug policy, prostitution, and any given internet list of dumb laws.<sup>2</sup>

Thus our starting point is that, by definition, all harms are bad; it is always wrong to harm others. If this seems trivial to you, just wait. As we will see, it will mean that some things which look like harms actually are not — in some cases, you can kill someone without harming them.

### **3.2.2 Definition of harm**

With that background in place, what is the relevant notion of harm? For the view loosely based on Thompson and Feinberg, I will say that

x is a harm if and only iff x is a wrongful setback of interest.

That means there are two parts to a harm. There has to be (1) a setback of interest and (2) it has to be done in a wrongful way. This give us a framework for understanding when a harm is present. We first identify the setback. Then we ask whether it was wrong. If we can do both, there was a harm. If we can do one but not the other, there is no harm.

But before we can do that, we need to get clear on what counts as a setback of interest and the relevant notion of wrongfulness.

#### **3.2.2.1 Setback of interest**

Let's start with the notion of a setback of interest. Obviously, we first need to know what interests are.

##### **3.2.2.1.1 Interests**

An interest in this sense is anything you have a stake in. If you invest in a company, you now have an interest in it. There are a bunch of ways of thinking about this, but I'm just going to say that:

---

2. Note that this doesn't automatically show that any of these things shouldn't be illegal. Drug use often negatively affects those around the user; prostitution is often bad for the prostitutes. It just reframes the question to whether these are harms in the relevant sense and gives us a means of answering that question.

You have an interest in  $x$  only if changes in  $x$  can affect your well-being.

Notice I didn't say something like 'the changes can affect you in a negative/positive way'. That's because of the harm principle lurking in the background. Remember, if something turns out to be a harm, we are allowed to criminalize it. If I give you a cookie, that positively affects you. If the temperature of the room is a bit colder than you'd like, that negatively affects you. But if we say that the room being a bit chilly is a harm, given the harm principle, that means we may be okay with using the full might of the criminal law to enforce HVAC settings.

### **3.2.2.1.1.1 Well-being**

Well-being is usually something which is more spread out over time. It's the sort of thing we have in mind when we talk about living a happy life. A happy life doesn't mean that you were happy at every single moment. The tricky thing with thinking in terms of well-being is to pick the right time interval for assessing it.

You don't want it to be too long. You don't want it to be assessed over your entire life — the sort of thing you consider on your deathbed in old age when you think back over your entire life and wonder whether it was worth anything. That's because lots of important events will get lost. That breakup in college was really painful and tough. But from the perspective of your entire life, it probably won't even register.

It also better not be so short that having an unusually bad cup of coffee from the student center on Tuesday morning damages your well-being. The idea is that when we think about how things are going for us, all the minor good and minor bad should normally drop out.

Thus for our purposes, I'm going to stipulate (without argument) that the right time-frame for assessing well-being is about a week. That is, things affect your well being when they affect the quality of your week. Imagine that on Friday you grab a drink with a friend who you haven't seen since last Friday. When you ask 'how have you been?' You are expecting to hear about the great date they had, the exam that they aced, the major fight they got into at work. If they start going on about how the cup of coffee they had Tuesday morning was terrible, you'd be pretty surprised or confused.<sup>3</sup> That sort of

---

3. Assuming that your friend is not a serious coffee snob. I had a good friend in graduate school who took his coffee super seriously (he taught me how to roast green coffee beans). I could imagine one bad coffee really sticking in his craw so much that it would still be on his mind several days later.



thing just usually isn't a big enough deal to actually affect how good or bad their week was; it doesn't usually affect their well-being.

### **3.2.2.1.2 Setbacks**

We now know that the relevant interests are the sorts of things which can affect the quality of your week. Now we need to talk about when an interest is set back. I know. I can see your eyes rolling already. Do we philosophers really have to overanalyze everything? Yes. It's what we do. This is actually interesting, I promise....

I'll start with the definition I'm headed towards, then tell you why we need it:

x setback a person's interests if and only if she is worse off than she was or would've been if x hadn't occurred.

Why not just say 'worse off than she was'? That covers most cases. You steal my bike, I am worse off than I was before you stole it.

The answer is death. That's right, death. Ok, not just death. There are other cases. If I steal your winning lottery ticket and replace it with a loser before you realize it is a winner, so you are no worse off than before I stole it. Adding the counterfactual condition 'or would've been' allows us to still say the theft was a harm.

If you'd like to join me for a brief, though relevant, tangent about death and its relevance for the counterfactual, read on, otherwise, you can skip ahead.

### **3.2.2.1.2.1 Death**

We are now to one of my absolute favorite puzzles in all of philosophy: how can your death be bad for you?

This goes all the way back to Epicures in ancient Greece, who used this to argue that it is irrational to fear death. But I'm getting ahead of myself. Here's the puzzle, paraphrasing Epicures

Where I am, death is not; where death is, I am not.

Do you see it? Let's walk through the puzzle. When you are dead, you stop existing. There is no more you.<sup>4</sup> If you still exist, you are only mostly dead (send your friends to

---

4. NB, if you are religious and believe in an afterlife, then in this sense, you believe that no one ever really dies. (C.f., Pharrell's project <http://nooneeverreallydies.com/>). You would agree with Epicures that it is irrational to fear death; it might of course be rational to fear what happens after death.

Miracle Max, ASAP!<sup>5</sup>). So, when you are dead, there is no one there for death to be bad for. Therefore, death cannot be bad for the person who dies.

Reactions to this puzzle differ. Epicures thought it (along with some other arguments) was liberating —it shows it is irrational to fear death. Others, myself included, see it as posing a problem to be solved. You might take a middle ground and say something like, yeah, it's irrational to fear death itself, but perfectly rational to fear dying. However, I think Nagel was right when he wrote "I think I should not fear dying if it weren't followed by death"<sup>6</sup>

### **3.2.2.1.2.2 Other problems**

There are other problems for this whole picture of setbacks of interest. Seana Shiffrin, for example, gives the example of a rich asshole whose hobby is going around and dropping bricks of gold on people's toes and then running off. She calls this jerkwad Richie Rich.

If you've never broken a toe, consider yourself lucky. I've had more than my share (thanks, judo). It is a particular kind of misery, because you can still basically walk and do stuff. But you live in constant fear bumping into furniture, stepping wrong on stairs, having an ant crawl over your sock, et cetera. It sucks.<sup>7</sup>

However, at least for me, the misery of a smashed toe would be substantially offset by the \$500k which came to me in the form of the shiny new 400oz gold bar resting on my foot. My going rate for allowing you to smash a (non-big) toe is, conveniently, about \$250k.

Thus if Richie Rich drops the bar on my toe, I am not worse off than I was. The bar is worth more than my smashed-toe-rate. And I'm also not worse off than I would've been if he had dropped the bar on someone else's toe. Sure I could walk without wincing. But I'd much rather limp with a half a million. So, on this picture, there would be no setback of interests and therefore no harm.

---

5. Then have fun storming the castle.

6. [reference] In the very unlikely chance that you were wondering what my favorite footnote in all of philosophy is, now you know.

7. The terror of coughing /sneezing when you have broken ribs during cold season is far worse. Just FYI

### **3.2.2.2 Wrongfulness**

The second component of harm is that it has to be wrongful. That makes sense when we remember that we are starting from the assumption that every harm is something that we can punish or coercively prevent people from inflicting.

There are many ways to make sense of the wrongfulness criteria. For the most part, any major moral theory can be slotted in here. However, we are going to follow Thomson and Feinberg in understanding the wrongfulness component as a rights violation. Therefore, a harm is a rights violation which makes the victim worse off than they were or would've been.

### **3.2.2.3 Odd cases**

It's worth emphasizing that, on our definition, to have a harm there has to be both the setback and the rights violation. If you only have one but not the other, there is no harm in the relevant sense.

In many cases this is straightforward. If you steal my bike, I am worse off than I were before (now I have to drive; I have to buy a new bike). The theft violated my property right to control what happens to and with my stuff. Thus you've harmed me. Easy peasy.

However, in some cases, applying this concept of harm leads to some counterintuitive results. Let's work through a couple of cases; doing so will let us see how the view works in action.

#### **3.2.2.3.1 Harmless rights violations**

Suppose you are sitting at a coffee shop and have just finished reading a newspaper.<sup>8</sup> When your back is turned, I sneak up and steal your newspaper. Is this a harm?

Well, it seems like I violated one of your property rights. You bought the newspaper. Sure, you were intending to toss it in the recycling or leave it for the next patron. But I took it before you did that. So, it looks like what I did was wrongful.

Are you worse off than you were or would've been? Probably not. In fact, I might of just saved you a trip all the way to the recycling bin, so you might be better off. Of course,

---

8. Newspapers were these things that everyone used to read. Kind of like the internet but on paper. Just bear with me.....

we might think up ways in which there was a setback —perhaps you now feel unsafe or your confidence in humanity has been shaken to the core. But at the very least, only some people are going to respond that way. And for most people these won't be setbacks that rise to the level of affecting your well-being. Therefore, it looks like there was no harm.

### **3.2.2.3.2 Hurting without harming**

Suppose that Scarlet kills Violet in genuine self-defense. Does Scarlet harm Violet?

Actually, before we can get to that, I need to clarify what I mean by genuine self-defense since the self-defense defense has been horrifically perverted by legislation in places like Florida.<sup>9</sup> Genuine self-defense means that the defender faced a lethal threat (they were being attacked by someone trying to kill them), the means were proportional (you can't shoot someone threatening to punch you), and had no escape (some jurisdictions carve out the 'castle doctrine' such that you have no duty to retreat in your own home).

Okay, back to the case at hand. Does Scarlet harm Violet? That may seem obvious. She kills her. How could that not be a harm? But remember, the question is whether she harms her in our sense; whether she wrongfully setback some of Violets' interests. So, let's work through it.

Does Scarlet setback Violet's interests? Yep. Which ones? Plausibly all of them (ignoring the complication around death mentioned above). It's sometimes useful to summarize this by saying that Scarlet hurts Violet.

Does Scarlet violate one or more of Violet's rights? If so, which ones? This gets complicated and there are several possibilities. For our purposes, I will just mention two ways this can go.

If the relevant right is a right to life, then it looks like there has been a rights violation. However, that means there has been a harm. Given the harm principle, that means we

---

9. There are plenty of legitimate arguments we can have over exactly what should count as self-defense. My philosophy of law class spends literally half the semester working through them; and I don't want to pretend I know what the answers are. However, I will say the subjective perceptions of the killer cannot be determinative in anything like a just society. (IIIRC, the NRA and other backers of 'stand your ground' laws agree and characterize the Florida law/cases as mistakes)

would be justified in punishing people who kill in genuine self-defense.<sup>10</sup> Violet teams up with the state to force Scarlet to choose between dying and going to jail. That does not seem plausible.

One common response is to invoke what's called the forfeiture theory. When someone unjustly tries to kill another person, they (temporarily) give up their right to life during the attack. Thus the attacker does not have a right to life that could be violated. Therefore, when Scarlet kills Violet, there's no harm.

There are several problems with the forfeiture theory. It creates complications with third-party assistance (whether the attacker forfeits their right overall or just to the victim affects how we deal with cases where a third person mistakenly thinks someone is being attacked and kills the attacker on their behalf). It also leads to this weird picture of rights. Remember that rights are supposed to be absolute prohibitions on people doing things to you. To see this, consider a wrinkle in how self-defense works:<sup>11</sup> Suppose that when Scarlet pulls out a gun, Violet stops trying to attack. But Scarlet is all pissed off and intends to kill Violet anyway, now Violet can kill Scarlet in genuine self-defense. If we are thinking about what's going on in terms of the forfeiture theory of rights, it all looks weird: first Violet has no absolute protection against being killed (which Scarlet has), then Scarlet's right disappears and Violet's reappears. Absolute no longer means absolute.

An alternative is just to be more careful about what the relevant right actually is. You could say that a 'right to life' is just a convenient shorthand. What we really have is a right not to be killed unjustly. If the relevant right is a right to not be killed unjustly, then we ask when is it justifiable to kill. The answer will include what I said above in describing genuine self-defense. In other words, on this version of the right, self-defense is one of the built-in exceptions to the prohibition on killing. Thus on this picture too, there will be no rights violation when Scarlet kills Violet and therefore no harm.

### **3.3 Informational harms -- setbacks**

Let's start by considering some of the ways personal data can be used to setback people's interests. We'll then turn to whether these setbacks involve rights violations.

---

10. Actually, this doesn't have to be true. There are other kinds of defenses under the criminal law that could get her off the hook. But that's too far afield.

11. Yes, there actually have been cases like this.

### **3.3.1 Exploiting weaknesses**

We all have bad habits and other psychological weaknesses. In many cases, we expend great effort keeping others from knowing them. The fact that these habits and vulnerabilities may be present in personal data about us or inferred from such data creates a significant opportunity for others to exploit them in ways which may significantly set back our interests.

Barocas and Nissenbaum point to some dimensions of this in discussing the way that targeted advertising online

might not only lock individuals into past habitual choices from which they would like to escape, but may open them to manipulation and illegitimate control by others. If someone can identify your weaknesses and vulnerabilities by closely monitoring past behaviors and dispositions, that person may be able to shape your choices, actions, transactions, and purchasing decisions in ways that do not accord with principles and purposes to which you are committed. Even if you succeed, in your deliberate actions, to stay true to these purposes and principles, others may have their own reasons for targeting your weaknesses, prejudices, or vulnerabilities, and, thereby undermining your autonomy. {Barocas:2009ws}

Now, it will matter how significant these weaknesses are and what behaviors they can get you to do. Remember, for something to be a harm, the setback has to affect your well-being. Take a minor bad habits like spending more time than you'd like on sites dedicated to celebrity gossip. If an online advertiser cunningly lures you back to such sites, that does set back your interests. But if it's just a few more minutes a day, the setback may not rise to the appropriate level for it to genuinely count as a harm.

That said, we should be sensitive to cumulative effects. We can imagine cases in which each manipulative advertiser succeeds in stealing a few minutes a week from you. On its own, not enough for a harm. But together, you are wasting a couple of hours a week on sites which you don't really want to be visiting. That would potentially be a harm. (If you were stuck for 2 hours at the DMV last week, your friend asking you about your week would not be surprised to hear about it).

### **3.3.2 Stalking and worse**

It should not be a surprise that the easy availability of personal location data enables stalking and other execrable behavior. After all, the actress Rebecca Schaeffer was murdered in 1989 by a killer who got her address from the DMV.<sup>12</sup>

Home addresses are fairly easy to find online. There probably aren't too many companies who can access a given person's current location. Though historical data about a person's movements is normally just as good; we are far more creatures of habit than we realize.

How do they get it? There is a bit of facial recognition floating around out there, but it's not usually a source of location data. Usually, it's our devices.

#### Car

Researchers have identified bluetooth signal harvesting devices on roads.<sup>13</sup> [Ref] Automated License Plate Readers are heavily used.<sup>14</sup> The wheels on modern cars report tire pressure via uniquely identifiable bluetooth signals.

Obviously, your cell carrier always knows where your phone is located. As long as 3 towers can get your phone's signal, your location can be pretty exactly determined. If your phone pairs with public wifi signals, your movements can be tracked with a great deal of precision, likely even your position inside a store, by whoever is operating the service.

Many of us frequently divulge our current location through social media posts. And, even when we do not directly post to tell the world that we are at the Starbucks on Reseda right now, we often reveal locations accidentally. Many phones and cameras automatically tag the picture with the GPS coordinates at which it was taken. Some platforms like Facebook now remove much of this metadata when photos are uploaded. But this is not universally the case.

---

12. See *Margan v. Niles*, 250 F. Supp. 2d 63, 68 (N.D.N.Y. 2003). Passage of the Driver's Privacy Protection Act (DPPA), 18 U.S.C. §§ 2721–2725 (2000), followed shortly thereafter in 1994. *Margan*, 250 F. Supp. 2d at 68–69.

[Refs from Nissenbaum]

13. [https://link.springer.com/chapter/10.1007/978-3-662-45317-9\\_3](https://link.springer.com/chapter/10.1007/978-3-662-45317-9_3)

[https://www.schneier.com/blog/archives/2008/04/tracking\\_vehicle.html](https://www.schneier.com/blog/archives/2008/04/tracking_vehicle.html)

14. <https://www.eff.org/pages/automated-license-plate-readers-alpr>

Indeed, given tools like Google's reverse image search, it is fairly trivial to figure out where a picture was taken.<sup>15</sup> You upload or provide a link to the image and the service finds other pictures likely taken in the same location. Check those out and you'll likely find someone who has included metadata or helpfully explained where it was taken.

[ToDo]

- 'Spouseware'

- Strava run data gives away locations of US military bases

<https://www.theguardian.com/world/2018/jan/28/fitness-tracking-app-gives-away-location-of-secret-us-army-bases>

- Japanese actress attacked after hi-resolution photo identified her location

A man arrested on suspicion of stalking a female pop idol used the reflections of her pupils in photos she shared on social media and Google Street View to find where she lived

<https://www.usatoday.com/story/news/world/2019/10/11/japan-man-arrested-stalking-pop-star-using-photos-eyes-pupils/3942667002/>

### **3.3.3 Threats to autonomy**

Hopefully, at this point in class, you're sensitive enough to concerns about autonomy, that the role privacy plays in providing the conditions for autonomy leaps fairly quickly to mind. Indeed, this connection lurks in the background for several other ways invasions of privacy can set back your interests.

Nissenbaum helpfully summarizes some of these connections

freedom from scrutiny and zones of "relative insularity" are necessary conditions for formulating goals, values, conceptions of self, and principles of action because they provide venues in which people are free to experiment, act, and decide without giving account to others or being fearful of retribution. Uninhibited by what others might say, how they will react, and how they will judge, unhindered by the constraints and expectations of tradition and convention, people are freer to formulate for themselves the reasons behind significant life choices, preferences, and commitments..... autonomy touches many dimensions of peoples' lives, including tastes, behaviors, beliefs, preferences, moral commitments, associations, decisions, and choices that define who we are. [From nissenbaum 2004 p.130] a

---

15. [https://support.google.com/websearch/answer/1325808?p=ws\\_images\\_searchbyimagetooltip&visit\\_id=637096219113096997-2863741639&rd=1](https://support.google.com/websearch/answer/1325808?p=ws_images_searchbyimagetooltip&visit_id=637096219113096997-2863741639&rd=1)



Thus invasions of privacy can set back a person's interests in being autonomous. Those are deeply important interests indeed since they underpin all sorts of other things which are important.

### **3.3.4 Relationships**

One theme throughout our discussion of personal data concerns how the information we share with others is deeply intertwined with our relationship to them. Nissenbaum summarizes some of these connections

Information is a key factor in the relationships we have and form with others.... controlling who has access to personal information about ourselves is a necessary condition for friendship, intimacy, and trust....[D]istinctive relationships, for example individual to spouse, boss, friend, colleague, priest, teacher, therapist, hairdresser, and so on, are partially defined by distinctive patterns of information [From nissenbaum 2004 p.130]

Thus our interests in maintaining the integrity of our relationships may be set back by violations of privacy.

### **3.3.5 Democracy**

Insofar as democracy is valuable and people have interests tied up with living in a political system which is based in and responsive to the needs of its citizens, invasions of privacy can set back these interests. Nissenbaum summarizes some of these concerns

privacy is essential to nourishing and promoting the values of a liberal, democratic, political, and social order by arguing that the vitality of democracy depends not...on the concrete protection against public scrutiny of certain spheres of decision-making, including but not limited to the voting booth. Privacy is a necessary condition for construction of... "social personae," which serves not only to alleviate complex role demands on individuals, but to facilitate a smoother transactional space for the many routine interactions that contribute to social welfare. Similar arguments...[defend] robust protections of medical information on grounds that individuals would then be more likely both to seek medical care and agree to participate in medical research. In turn, this would improve overall public health as well as social welfare [From nissenbaum 2004 p.132]

### **3.4 Informational harms -- rights violations**

Now that we have a solid definition of harm, we can work through the cases we are concerned with to try to determine when people are actually harmed.

Remember, on our notion of harm, there can be bad things which happen to people which do not count as harms. There's both a minimal threshold of severity covering the setback of interests and the requirement that we be able to locate a right that's violated. Therefore for each alleged harm from informational privacy violation, we will ask what interests are setback and which right is violated.

Thus before we jump in, it will be worth doing the thing I've put off until now: asking what rights there are around privacy.

### **3.4.1 What kind of rights?**

#### **3.4.1.1 Non privacy rights**

The the rights violated might not have anything to do directly with informational privacy.

When your personal data is used to fraudulently or coercively take away your money or your property, this looks like a straightforward violation of your property rights. For example, if I pay you off to prevent you from revealing my Netflix history to embarrass me in front of my pretentious art-film-snob friends, you have violated my property right to the blackmail money.

Rights to security of person can be compromised when personal data is used to enable stalking, intimidation, or other crimes.

#### **3.4.1.2 Right to protected sphere**

But in some cases, we won't be able to find a more commonplace rights violation. For example, the potential harms related to surveillance will not likely violate any property right or threaten your personal security. Thus we will need to look for rights which have to do directly with privacy.

If the idea of privacy rights strikes you as implausible —some writers maintain that all alleged privacy rights are derivative from more familiar rights—it's worth noting that in American law we do recognize legal privacy rights on their own. Under the common law, there is commonly (hah!) held to be a right to informational privacy. This is sometimes referred to as *tort privacy*. There is a tort of privacy invasion which applies

when the defendant has intruded 'into [plaintiff's] private affairs' or '[p]ublic disclosure of embarrassing private facts about the plaintiff.'<sup>16</sup>

There is a ton of controversy over the details of such rights and over what they covers. We will talk about what 3 families of view which attempt to explain privacy rights and what they cover in what follows.

Before we jump in, it's worth noting that we can characterize all of them as entailing the existence of a *protected sphere of information*. We see this all the way back in the famous paper which kicked off the discussion of privacy. Warren and Brandeis were explicitly concerned with protecting information about the 'private life, habits, acts, and relations of an individual'<sup>17</sup> This means that people have a right to determine how and when information within that sphere is communicated to others.<sup>18</sup> For our purposes, trespassing into that sphere will be wrongful. Thus if information within that sphere causes a setback of interests, there will be a harm.

### **3.4.2 Accounts of right to protected sphere**

Let's set aside the difficult of identifying the precise rights at stake. Pretty much any view is going to suppose a protected sphere of information covered by a right to privacy. We can make some progress by considering how to justify the necessity of such a sphere.

We will explore 3 competing accounts: accounts based in negative liberty —non-interference—concerns, accounts based in republican concerns about non-domination, and accounts that expand beyond the individual to consider concerns around relationships.

#### **3.4.2.1 Negative Liberty account**

Negative liberty accounts of the right to a protected sphere of information will be concerned with the role protecting certain information plays in being free.

---

16. William L. Prosser, Privacy, 48 CAL. L. REV. 383, 389 (1960).

17. Warren & Brandeis, 216.

18. Muller, p.14

Why is this called 'negative liberty'? The distinction between negative and positive conceptions of liberty goes back to a famous paper by Isaiah Berlin. Very roughly, if we ask 'When is a person free?', on the negative liberty account we get the answer: When others do not interfere with her choices. On the positive liberty account, we get: When she is good.

I will sometimes refer to this account as the 'liberal' account. This is a standard category of view in political philosophy. It is not necessarily 'liberal' as the term gets used in American political discourse. Where there are connections between this sort of liberal political philosophy, they exist because some members of the American Democratic party are attracted to these views; not the other way around.

Let's spell out this picture in more detail and see how it fits with the idea of a protected sphere of information. Since the negative liberty picture likely strikes most American readers as obvious, we'll start with positive liberty to make it easier to see what the negative liberty picture implies.

#### **3.4.2.1.1 Positive liberty**

On the positive liberty account, the idea of freedom is somewhat perfectionist<sup>19</sup>. To be truly free, you need to have certain kinds of personal development and probably certain beliefs.

One place, indeed the original place for Berlin, this can be found is a broadly Marxist picture. For Marx, in a capitalist society workers are told by all sorts of social forces that the owners of the factory are supposed to be the owners and they are supposed to be just a cog; their labor belongs to those who pay for it. The Marxist thought is that you have to be freed from this false conception of the world before you can be truly free. Something like this is built into the positive liberty conception.

Alternatively, suppose you grow up in a society with a heavy caste system, the lives open to you are circumscribed —you can be a tanner, a plumber, or a butcher, those are your only choices. You can choose freely among them; no one will make you be a butcher when you want to be a plumber. But on this sort of picture you aren't truly free

---

19. Perfectionism in ethics is concerned with things like achievement and excellence. For example, on other views, an athlete setting a world record might be praiseworthy because it makes fans excited. But the perfectionist will claim that achieving someone no one else ever has is in itself valuable.

unless you recognize that the society has trapped you into those choices. There are other careers that you could've pursued — actor, lawyer, fishmonger— but for the caste system.

That said, I mention this mainly for completeness and to illustrate the negative account. We won't talk much about positive liberty in relation to privacy. Though it should be clear that there's not going to be much presumption in favor of privacy, for those who haven't been properly situated / trained and don't have the conditions of positive liberty.

#### **3.4.2.1.2 Negative liberty**

The negative liberty view on the other hand doesn't require anything of the individual's beliefs or self-conception. It just requires that others don't prevent her from doing things that she wants to do.

To be clear, this isn't saying that all restrictions on action are wrong. Normally, a negative liberty approach will get fleshed out to specify that some interferences, e.g., restrictions that prevent people from murdering children do not diminish freedom. In other classes when I talk about these views (and sometimes in our class when talking about autonomy), I usually describe them as something like:

You are free when you are able to live the sort of life that seems worth living to you compatible with equal liberty for all.

That's roughly in line with this tradition.

#### **3.4.2.1.3 Right to protected sphere**

Thinking about privacy from the negative liberty perspective puts an emphasis on the ability to exclude others from access to certain kinds of information because such access interferes with your ability to live the life you want. To be free, you must be able to, among other things,

- Keep other people from interfering with your choices
- Preserve available options for your choices
- Choose in accordance with your genuine preferences
- Retain control over how you present yourself to others
- Control much others know about our private lives

In short, any time your ability to make choices on the basis of what you genuinely care about is restricted, you are less free. Assuming that it is always better to be more rather than less free, that means such interferences are worse for you (though obviously things come in degrees).

The interference doesn't have to be completely obvious, like someone blocking the door you want to enter. Indeed, the interference can turn heavily on your own self-censorship. If you know that all of your internet browsing history will be publicly available, you probably will not use the internet the way you would if others weren't watching. You might not watch the frivolous videos that help you relax after work. You might not google for embarrassing health information. You might not use the internet to help you explore your sexual identity. All of these would be significant interferences with your ability to live the kind of life you want to live, though no one else actively and intentionally prevented you from doing these.

For understanding how this view differs from the others, let me emphasize that the focus here is on actual non-interference. If people could interfere with your choices but do not, there's no problem for your liberty. Obviously some caveats will apply. If your knowledge of their ability to interfere inhibits you from acting, you're still having your liberty restricted on this view.

Consider our example of the Data-Hoarder. Given the assumption that the database is perfectly secure and that she will never have the urge to look at your data within it, her collecting the data doesn't seem to interfere with your choices. Thus the emphasis on actual interference for the negative liberty account seems to lack the resources to say that her collecting and storing all your information actually trespasses into your protected sphere of information. It's only if she looks, that your right to a protected sphere of information would be violated.

#### **3.4.2.1.4 Objections**

[ToDo]

- Cannot explain intrinsic wrongness of ubiquitous surveillance. That's because what's wrong is the interference, not the data collection itself. If government (etc) is appropriately constrained in what they do with your stuff, it's hard to motivate a problem [See Stahl]
- In general, as long as the data is not used to interfere with individual liberties, hard to explain what's worrisome [See Stahl]
- How concerned we should be about collection of some data doesn't depend on what the data is per se, it depends on how likely it is that it will be used to interfere

#### **3.4.2.2 Republican account**

Where the negative liberty view starts with the danger interference with one's choices poses for autonomy, the neo-republican view focuses on the existence of certain power relations. For the neo-republican, the mere existence of a domination relationship is objectionable. Where the negative liberty view has trouble explaining what is problematic in cases where no actual interference is likely, the neo-republican can complain that as long as the option exists, it is unacceptable because people are not on an equal footing.

By the way, note that this is not 'republican' in the sense of the American Republican party. There are some ideological connections, but where they exist, it's because some members of the Republican party are attracted to this position in political philosophy, not the other way around.

Let's understand domination as:

A is dominated by B if B has the option of interfering with A's choices, regardless of whether the option is taken or not

Thus for the neo-republican, A is free when no one dominates her.

This provides in some ways a more pure conception of freedom — where the negative liberty view is concerned with actual interference, this view gives grounds for complaint whenever it is possible that someone interferes with her choices.

The fact that a government or other power is in position to use people's information constitutes domination. Thus it looks like there are more resources on this view to complain about violations of privacy. Indeed, a serious strength of this picture is that concerns about invasion of protected sphere of information need not track probability of interference.

#### **3.4.2.2.1 Right to protected sphere**

Thus we can see that on the neo-republican account, the right to a protected sphere of information is grounded in concerns about preventing the existence of domination relationships. Information will fall within the protected sphere if the possession or use of information would impose a domination relationship.

#### **3.4.2.2.2 Objections**

[ToDo]

- A good regulatory / constitutional system could block information holders from using it. Thus the option is taken away and there is no domination. [See Stahl]

- Gathering information per se does not involve domination. Thus still may not have reason to worry about ubiquitous surveillance

### **3.4.2.3 Relational account**

Where the liberal and republican accounts have been focused on the individual and the effects of intrusions (or the possibility of intrusions) into the protected sphere on individuals, relational views are concerned that this obscures important effects of privacy.

On these views, privacy is a necessary condition of being able to form / conduct intimate (and other relationships).

For example, personal and professional relationships are governed by norms concerning what information is shared.

#### **3.4.2.3.1 Right to protected sphere**

[ToDo]

Relationships

Role of information sharing in relationships

Groups

Within groups, having certain shared secrets is important for group cohesion/bonding. Think of shared lore and history in sororities or fraternities.

Workplace

Think about information sharing in the workplace. Since you often need to work productively with people who have very different political or religious views, norms of professionalism normally require that people do not go into too much detail about these topics.

Similarly, organizational coherence often requires those in leadership to maintain some secrets from those below them. More importantly, and I've seen this a gazillion times, getting people to agree and be on the same page with a controversial decision often means allowing them to agree for different reasons. If all the differing reasons behind a controversial decision were public, it would be very difficult to convince those outside the decision making group to go along.



Consciousness raising

- Marginalized groups need privacy from dominant groups in order to recognize the legitimacy of their experiences, deal with internalized oppression, and prepare political action

Civil inattention

- Nagel: Let strangers be strangers; necessary for civil society

### **3.4.2.3.2 Objections**

[ToDo]

- Stahl: Most relationships not threatened per se by surveillance
- Still focused on control of information which individuals want / need to keep private

## **3.5 Can machines violate privacy?**

Let's close off our discussion of privacy-related harms by considering a question whose relevance will become clearer once we start discussing responsibility. Can a machine violate privacy?

For example, the original business model of gmail was to serve users ads based on keywords located in emails they send and receive. There was no Google employee behind the scenes reading your emails and determining which ads to show you. The process was entirely algorithmic. An outdoor brand would pay Google to advertise to people interested in, say, hiking. Google would look for words like 'tent' or 'hiking boots' in emails and serve ads to people sending/receiving those emails. To people who felt this violated their privacy, the response (I'm paraphrasing) would be 'but you and the recipient are the only humans looking at the emails. How could that violate privacy?'

Let's try to think through whether there could be harms in these sorts of cases using the various theories we've discussed.

### **3.5.1 No. Understanding matters**

Before we get to that, let's mention one very flat-footed reply: No machine can violate privacy because machines do not understand what they are 'reading'. The argument thus may go:

- (1) A violates B's privacy regarding x only if A understands x.

**Very rough draft: Do not circulate**

(2) Machines cannot understand anything.

Therefore, (3) Machines cannot violate privacy.

We could fight about premise (2). But that would take us on a long detour through philosophy of mind and philosophy of language. So we won't.

More importantly, we don't need to worry about (2) unless (1) is true since it makes the privacy violation depend on the existence of understanding. If (1) is false, then it may be that privacy can be violated even when there is no understanding.

Why would we believe (1)? If you start with the thought that what we care about is other people knowing stuff about us, it will seem plausible. Suppose that someone steals your bag containing the diary in which you've written all your deepest darkest secrets while you are traveling in a country where no one speaks English (and phones which can translate text in a picture don't exist). Are you worried about your privacy? If you are, this view would claim, it's only because you're worrying that someone who speaks English will come along and read it. If you are certain that won't ever happen, then you shouldn't be concerned. And, certainly, we wouldn't say that your privacy has already been violated as soon as the theft occurred.

But we've already seen that the harms we're concerned with do not turn on the question of whether anyone/anything understand what they read. On all of our accounts, the harm can arise from self-censorship. And, thinking of cases like the Argument Tweeter or People-Rating, should remind us that we care about the ways our information can be used regardless of whether there is an understanding agent in the loop.

### **3.5.2 Application of 3 approaches**

[ToDo]

Liberal

Neo-republican

Relational